Many Uses, Many Annotations for Large Speech Corpora: Switchboard and TDT as Case Studies

David Graff and Steven Bird

Linguistic Data Consortium
University of Pennsylvania
3615 Market Street, Philadelphia, PA 19104, USA
`http://www.ldc.upenn.edu`

Abstract

This paper discusses the challenges that arise when large speech corpora receive an ever-broadening range of diverse and distinct annotations. Two case studies of this process are presented: the Switchboard Corpus of telephone conversations and the TDT2 corpus of broadcast news. Switchboard has undergone two independent transcriptions and various types of additional annotation, all carried out as separate projects that were dispersed both geographically and chronologically. The TDT2 corpus has also received a variety of annotations, but all directly created or managed by a core group. In both cases, issues arise involving the propagation of repairs, consistency of references, and the ability to integrate annotations having different formats and levels of detail. We describe a general framework whereby these issues can be addressed successfully.

1. Introduction

Any well-constructed corpus of speech data can provide a valuable resource for a wide variety of uses in language research and technology development, especially if the corpus is centered on common and naturally-occurring speech events. Both the potential and the motivation for re-use increase with the size of the corpus: larger corpora provide a better representation of linguistic diversity and variability, and so are richer objects of study for any research goal; also, the expense and effort that go into the creation of a large corpus, typically on behalf of a particular research program, can provide powerful leverage for researchers involved in other projects or areas of study.

It is often the case that a new use of a corpus will require a new annotation that was not part of the initial corpus creation effort. But for large corpora, it often happens that the new annotations can be applied only to subsets of the collection, depending on the difficulty of the new task and the resources available to carry it out. In this case, it becomes increasingly important, and increasingly difficult, to maintain the coherence of shared resources.

The Linguistic Data Consortium (LDC) has been involved in managing multiple annotations of two large speech corpora: the Switchboard Corpus of conversational telephone speech (SWB), and the Topic Detection and Tracking corpus of broadcast news (TDT). In the case of SWB, the original corpus creation and all subsequent annotations have been done by others outside the LDC, and we have acted simply as the central point of contact and access for all users and annotators; in the case of TDT, LDC personnel have been responsible for most forms of annotation done so far.

The following sections will review these corpora in terms of their overall content, the particular types of annotations that have been applied to them so far, and the relative coverage of those annotations (including the overlap that exists among them). We then discuss some of the problems that arose due to localized data repairs that were applied in some annotations, and measure the extent of referential consistency among the diverse annotations. Finally, we describe how the application of annotation graph structure to the various derivative data sets provides a powerful and flexible means for comparing and integrating their contents.

2. The Switchboard Corpus

Switchboard-1, the first large collection of spontaneous conversational speech over the telephone, was collected in 1990 by Texas Instruments (TI) (Godfrey et al., 1992). The current published corpus comprises 2438 calls involving 520 native speakers of American English, recruited from all over the United States. The calls range between five and ten minutes in length, and each call contains a unique pairing of speakers. Individual speakers are designated by identification numbers, and information is provided as to their gender, age, education and dialect region. On average, each speaker appears in about nine calls (the range of calls per speaker varies between 1 and 32).

The speech data for each call is provided in the form of a two channel interleaved sample file; the two channels, designated "A" and "B", represent the mu-law encoded signal received from the telephone handsets of the two speakers participating in the call.

2.1. Initial transcription

TI created an initial set of manual transcripts, employing professionally trained transcribers equipped with analog tape copies of the sample files. These transcripts were then submitted to an automatic speech recognition (ASR) process to establish approximate time alignments at the word level.

The initial publication of the corpus in 1993 contained 2288 calls on 25 CD-ROMS (additional calls were being held in reserve at that time for future use as test data). A separate CD-ROM provided the transcripts in two forms (two files per call): (1) the original text files as created

by transcribers, including a generic header with information about the call, and speaker designations ("A" or "B") for each turn in the dialog, but no time annotations; (2) the time-aligned version of the text, with one word token per line, accompanied by the speaker designation, and the beginning time and duration of the word. The transcript release also included documentation and tables describing the corpus content.

Given the size, complexity and novelty of this effort, a variety of problems affected the first release of the corpus by the LDC:

- In about 200 calls, the transcribers had made mistakes in assigning the speaker labels ("A" or "B") to some or all of the turns in the transcripts. Because the speaker label is intended to represent the signal channel in the speech file, an error in the speaker label would direct the corpus user to the wrong channel when retrieving the audio data for a given turn in the dialog.

- Other more typical kinds of errors, involving misperceptions by transcribers about what was said, presumably affected nearly all files, but only to a very small degree.

- The ASR word alignment was applied to a single-channel (combined) version of the speech data, and so was not affected by speaker label errors, but it was relatively ineffective in regions where turns overlapped in time, and where the transcript contained omissions, additions or substitutions of words relative to what had been spoken.

- In about 30 calls, the two forms of transcript data did not match up in terms of the number of speaker turns present in the dialog.

- For about 30 additional calls, some or all of the transcript data were absent from the publication, because the text files were missing, incomplete or corrupted due to errors in preparing the publication.

In 1994, the LDC released an "interim" version of the transcripts, fixing all cases of the last two types of problems, and about one quarter of the files containing known speaker-label errors (particularly, those files where the label errors were limited to small portions of the transcript). The remaining speaker-label problems were documented as well as possible, but were not fixed, nor was any attempt made to correct other transcription or time-alignment errors.

2.2. Subsequent annotations

Since the initial release of SWB transcripts, a number of research sites have used this data, either with or without reference to the speech, as input to a range of divergent annotation projects. These are briefly described below, in roughly chronological order. Some of the resulting annotations are displayed in Figure 1.

Phrase-level time stamps (BBN)

The first application of SWB data for research in large vocabulary conversational speech recognition was in conjunction with the DARPA LVCSR project. In preparation for initial training on SWB, researchers at BBN created a modified version of the original TI transcripts, and circulated this among the LVCSR participants. The modification involved forming time-stamped "phrasal" regions from the word-level time-aligned transcripts, and assigning a unique identifier to each region. This produced a segmentation of the dialogs that differed from the turn units created by the original transcribers: a single time-stamped phrase might encompass parts of two consecutive turns, if the original transcriber had broken up the phrase in order to insert an interruption by the other speaker; also, multiple time-stamped phrases might be derived from a single turn, if that turn was considered too long for reliable ASR processing. The speaker-label errors mentioned above were not repaired in this process (though a list of affected files was circulated among LVCSR participants). This form of the transcripts has never been made available through the LDC.

Disfluency annotation (Penn Treebank Project)

As a preliminary step to prepare for Treebank annotation, 650 transcript files were selected from the LDC "interim" release and annotated for various types of disfluency that occurred in the spontaneous speech. This annotation was essentially text-based, relying entirely on the representation of disfluencies in the (non-time-aligned) transcripts (see Figure 1). The objective was to tag hesitations, stutters, word fragments, restarts, and a limited class of discourse markers, since these elements in the transcripts would constitute exceptions or barriers to syntactic analysis. These annotations are included in the current "Treebank 3" corpus, available from the LDC.

POS-tagging, parsing (Penn Treebank Project)

Building on the output of the disfluency annotations, the Penn Treebank Project applied part-of-speech tagging and syntactic parsing to the 650 files that had been selected (see Figure 1). In the course of this annotation, a small number of corrections were made to the text data, consisting mostly of repairs to punctuation and replacements of some incorrect words. Again, this annotation was done without reference to the audio data. It is currently available as the "Treebank 3" corpus.

Discourse annotation (Univ. of Colorado, SRI)

This annotation project set out to "model the speech act type of each utterance" and "model sociolinguistic facts about conversation structure..." (Jurafsky et al., 1997) It builds on the disfluency annotation mentioned above, using the same conceptual basis for segmenting dialog turns into phrases ("utterances") that are cohesive in terms of the speech acts being performed. An exhaustive segmentation of such utterances was carried out on 1155 conversations, and each utterance was categorized as to speech act or discourse function. All 650 files covered by the Treebank Project are included in this set. The data is available directly from the University of Colorado (Jurafsky, 1997).

Aligned Word	Part of Speech	Disfluency

```
B 19.44 0.16 Yeah,        ====================    B.22:  Yeah, / no one seems to be adopting it. /
B 19.60 0.10 no           [ SpeakerB22/SYM ]        Metric system, [ no one's very, + F uh,  no one wants ]
B 19.70 0.10 one          ./.                        it at all seems like. /
B 19.80 0.24 seems        ====================    A.23:  F Uh,  [ [ the, + the, ] + the ]
B 20.04 0.02 to                                      public is just very conservative that way in
B 20.06 0.12 be           Yeah/UH ,/,                refusing to change measurement systems,
B 20.18 0.50 adopting     [ no/DT one/NN ]           F uh,  money, dollar, coins, anything like that. /
B 20.68 0.16 it.          seems/VBZ to/TO         B.24:  Yeah <laughter>. /
B 21.86 0.26 Metric       be/VB adopting/VBG      A.25:  [ [ C And,  + C and, ] + C and ]
B 22.12 0.26 system,      [ it/PRP ] ./.             [ it + <breathing>,  it ] obviously makes no sense
B 22.38 0.18 no                                      that we're practically alone in the world [ in, + in ]
B 22.56 0.06 one's        [ Metric/JJ system/NN ]    using the old system. /
B 22.86 0.32 very,        ,/,
B 23.88 0.14 uh,          [ no/DT one/NN ]
B 24.02 0.16 no           's/BES very/RB ,/,
B 24.18 0.32 one          [ uh/UH ] ,/,
B 24.52 0.28 wants        [ no/DT one/NN ]
B 24.80 0.06 it           wants/VBZ
B 24.86 0.12 at           [ it/PRP ]
B 24.98 0.22 all          at/IN
B 25.66 0.22 seems        [ all/DT ]
B 25.88 0.22 like.        seems/VBZ like/IN ./.
A 28.44 0.28 Uh,          ====================
A 29.26 0.14 the,         [ SpeakerA23/SYM ]
A 29.48 0.14 the,         ./.
A 29.82 0.10 the          ====================
A 29.92 0.34 public
A 30.26 0.06 is           [ Uh/UH ] ,/,
A 30.32 0.22 just         [ the/DT ] ,/,
A 30.54 0.14 very         [ the/DT ] ,/,
A 30.68 0.68 conservative [ the/DT public/NN ]
A 31.36 0.18 that         is/VBZ just/RB very/RB
A 31.54 0.30 way          conservative/JJ that/DT
A 32.56 0.12 in           [ way/NN ]
A 32.74 0.64 refusing     in/IN refusing/VBG
A 33.60 0.12 to           to/TO change/VB
A 33.72 0.56 change       [ measurement/NN
A 34.94 0.48 measurement    systems/NNS ]
A 35.42 0.62 systems,     ,/,
A 36.08 0.26 uh,          [ uh/UH ] ,/,
A 37.04 0.38 money,       [ money/NN ] ,/,
A 37.62 0.30 dollar,      [ dollar/NN ] ,/,
A 37.92 0.46 coins,       [ coins/NNS ] ,/,
A 38.38 0.22 anything     [ anything/NN ]
A 38.60 0.18 like         like/IN
A 38.78 0.30 that.        [ that/DT ] ./.
B 39.34 0.10 Yeah
B    *    * [laughter].
A 40.96 0.04 And,         ====================
A 41.32 0.04 and,         [ SpeakerB24/SYM ]
A 42.28 0.36 and          ./.
A 42.88 0.20 it           ====================
A    *    * [breathing],
A 43.08 0.16 it           Yeah/UH ./.
A 43.48 0.46 obviously
A 43.94 0.22 makes        ====================
A 44.16 0.14 no           [ SpeakerA25/SYM ]
A 44.30 0.36 sense        ./.
A 44.66 0.06 that         ====================
A 44.72 0.12 we're
A 44.84 0.70 practically  And/CC ,/, and/CC ,/,
A 46.52 0.32 alone        and/CC
A 46.84 0.10 in           [ it/PRP ] ,/,
A 46.94 0.06 the          [ it/PRP ]
A 47.00 0.44 world        obviously/RB makes/VBZ
A 47.44 0.16 in,          [ no/DT sense/NN ]
A 48.52 0.04 in           that/IN
A 48.56 0.26 using        [ we/PRP ]
A 48.82 0.08 the          're/VBP practically/RB
A 48.90 0.22 old          alone/RB in/IN
A 49.12 0.40 system.      [ the/DT world/NN ]
                          in/IN ,/, in/IN
                          using/VBG
                          [ the/DT old/JJ
                            system/NN ]
                          ./.
```

Treebank

```
((CODE SpeakerB22 .))
((INTJ Yeah , E_S))
((S (NP-SBJ-1 no one)
    (VP seems
        (S (NP-SBJ *-1)
           (VP to (VP be (VP adopting (NP it)))))) . E_S))
((S (NP-TPC Metric system) ,
    (S-TPC-1 (EDITED (RM [)
                     (S (NP-SBJ no one)
                        (VP 's (ADJP-PRD-UNF very))) ,
                     (IP +)) (INTJ uh) ,
             (NP-SBJ no one)
             (VP wants (RS ]) (NP it) (ADVP at all)))
    (NP-SBJ *)
    (VP seems (SBAR like (S *T*-1)))) . E_S))
((CODE SpeakerA23 .))
((S (INTJ Uh) ,
    (EDITED (RM [)
            (EDITED (RM [) (NP-SBJ-UNF the) , (IP +))
            (NP-SBJ-UNF the) , (RS ]) (IP +))
    (NP-SBJ-1 the (RS ]) public)
    (VP is
        (ADVP just)
        (ADJP-PRD very conservative)
        (NP-MNR that way)
        (PP in
            (S-NOM (NP-SBJ-2 *-1)
                   (VP refusing
                       (S (NP-SBJ *-2)
                          (VP to
                              (VP change
                                  (NP (NP measurement systems) ,
                                      (INTJ uh) , (NP money) ,
                                      (NP dollar) , (NP coins) ,
                                      (NP (NP anything)
                                          (PP like
                                              (NP that)))))))))))) . E_S))
((CODE SpeakerB24 .))
((INTJ Yeah . E_S))
((CODE SpeakerA25 .))
((S (EDITED (RM [)
            (EDITED (RM [) And , (IP +)) and , (RS ]) (IP +)) and (RS ])
    (EDITED (RM [) (NP-SBJ it) (IP +) ,)
    (NP-SBJ (NP it)
            (SBAR *EXP*-1))
    (RS ])
    (ADVP obviously)
    (VP makes
        (NP no sense)
        (SBAR-1 that
                (S (NP-SBJ-2 we)
                   (VP 're
                       (ADVP practically) (ADJP-PRD alone)
                       (PP-LOC in (NP the world))
                       (EDITED (RM [) (PP-UNF in) , (IP +))
                       (PP in (RS ]))
                       (S-NOM (NP-SBJ *-2)
                              (VP using
                                  (NP the old system)))))))))) . E_S))
```

Figure 1: Multiple Annotations of the Switchboard Corpus

Phonetic transcription (ICSI)

The International Computer Science Institute (ICSI) at the University of California, Berkeley, began a project in 1996 to carry out fine-grained phonetic annotations on portions of SWB data (Greenberg, 1999). The selection of portions to transcribe, as well as the initial orthographic transcriptions, were apparently derived from the BBN phrase-level segmentation, and the resulting annotations are indexed by means of the unique identifier strings assigned to each phrase by BBN. In contrast to the other annotation projects, ICSI selected a sampling of 5100 phrases from a wide range of files, with each phrase ranging from 0.45 to 17.430 seconds (the majority are between 3 and 5 seconds). Two files received fairly exhaustive treatment for one side of the call, but for 1602 other files, the typical coverage of ICSI transcriptions is on the order of tens of seconds. Altogether, the combined annotations cover nearly 3.5 hours of speech. This data set is currently available from the Center for Language and Speech Processing at Johns Hopkins University (CLSP, 1996-7).

The initial stage of the project assigned time marks at the level of segmental boundaries, as well as syllable and word boundaries. In the second phase, time boundaries were applied at the word and syllable levels. Throughout the project, the labels assigned to phonetic segments (whether time marked individually or at the syllable level) were intended to accurately reflect the actual pronunciation in the signal, to a much finer level of detail than in previous corpora (e.g. TIMIT or the Boston University Radio corpus).

Complete resegmentation (ISIP)

In view of the importance of SWB as a multi-functional corpus, and the difficulties that have accompanied the original transcripts, an important effort was launched at the Institute for Signal and Information Processing (ISIP) at Mississippi State University, to conduct a complete review of the transcription data, applying a new speech transcription tool developed at ISIP with particular attention to locating known types of problems and avoiding known pitfalls in this type of annotation effort. By the time this project began, the SWB calls that had previously been held back as test data had been used over the course of several benchmark tests in the LVCSR project, and were now available for publication. The current release of SWB speech and transcript data now comprises 2438 calls. By agreement with the LDC, ISIP has made the complete set of SWB transcripts freely available on their web site (Mantha, 2000).

2.3. Rectification and Integration

With the completion of the ISIP review of SWB transcripts, it is now possible to assess the impact of transcription errors on the various divergent annotations. As an initial step to check for the magnitude of errors involving lexical content, we treated the LDC "interim" transcripts as a test set to be measured for error rates, using the ISIP transcripts as the reference text. Comparable versions of the two data sets were constructed by aligning the phrase-level ISIP time marks with the original TI word-level time marks, and the NIST scoring tool "SCLite" was used to calculate insertions, deletions and substitutions. The results are summarized in table 1. [1] It should be noted that word fragments and non-lexical tokens (e.g. "uh-hum" vs. "uh-huh") accounted for roughly 30% (over 19,000) of the substitution errors, and about 21% (nearly 34,000) of the insertion and deletion errors.

Units	Status	K	%	Per-file % range
phrases	correct	136	55.1	
	w/errors	111	44.9	7.6 - 90.9
words	correct	2895	94.8	77.4 - 99.5
	accuracy		92.8	
	all errors	220	7.2	0.8 - 27.9
	substit.	63	2.1	0.0 - 14.5
	deleted	95	3.1	0.3 - 15.5
	inserted	63	2.0	0.0 - 21.0

Table 1: Summary of word errors in LDC "interim" transcripts

The comparison of ICSI annotations to the other versions of SWB transcripts is somewhat more problematic. ICSI transcribers made corrections to the lexical content in accordance with their more detailed attention to the actual pronunciation of phrases, but they used different conventions regarding word hyphenation and disfluencies, and occasionally inserted annotations for non-linguistic events

(e.g. "breath", "mouthnoise", etc.) without the intended markup to distinguish these from lexical tokens. Still, we can estimate the upper bound on the number of corrections imposed by ICSI, again using the NIST SCLite scoring method. In this case, the LDC transcripts contained 96% of the words in the 3.5 hours of ICSI word-level data – there were, at most, 2% omissions and 2% substitutions in the LDC transcripts.

ICSI's use of the BBN phrasal time marks as the basis of phrase selection creates some additional difficulties:

- Correlation with original TI word-level time marks is imperfect at best; insertion errors predominated in scoring the LDC transcripts (6.2%, three times more than deletions or substitutions), yielding an overall word accuracy of 89.2%. This was due mostly to discrepancies at phrase boundaries.

- Some of the BBN "phrase boundaries" occur at impractical positions within syllables, causing some boundary tokens to be interpreted differently.

- Until a reliable word-level time marking is done for the ISIP transcripts, there will be no reasonable way to align the ICSI and ISIP annotations, due to significant differences in phrase segmentation.

Overall, these tabulations indicate that the lexical accuracy of the original TI transcripts was quite high, and the impact of word errors on downstream annotations, particularly the Treebank and discourse data, may be considered negligible.

A major challenge remains, however, in terms of integrating the various annotations. All three versions of time-marked transcripts (LDC "interim", ICSI and ISIP) assign unique identifiers to each turn or phrase unit, but each set has a distinct inventory of units that cannot, as yet, be cross-referenced to the other two sets in any reliable way. The only stable point of reference is the audio data, and the use of time offsets into the speech files.

3. The TDT Corpora

The design and content of the TDT corpora are described in other presentations at LREC-2000 (Wayne, 2000; Cieri et al., 2000). The present discussion will focus on the range of distinct annotations applied to the data, the relationships among them, and the problems involved in coordinating them.

3.1. Multiple data streams from audio and text sources

The audio recordings of broadcast (video and radio) sources in TDT were used to create a variety of textual data streams. For video sources (which have all been in English), the broadcast signal included closed-caption text data, which was converted to computer-readable form while the audio was being digitized; also, one video source (ABC News) provided full transcripts of its daily broadcasts to the public through a commercial transcription service. It is well known that closed-caption text tends to be incomplete relative to what is actually spoken during a broadcast, because

---

[1] LDC files that still contained speaker/channel labeling errors over some or all of their transcripts were not included in this scoring.

the maximum practical display rate, in words per minute, is slower than typical rates of speech. The full transcripts produced by commercial services, which are intended for use as a standard public record of broadcast content, are lexically correct to a high degree of accuracy, though they typically avoid the inclusion of any disfluencies – the text represents only what the speaker *intended* to say, and omits filled pauses, false starts, stutters, and the like. On the basis of this one source, then, it is possible to estimate a baseline of "word error rate" for closed-caption text, which can be useful when assessing other sources for which only the closed-captions are available.

The radio sources in TDT did not have publicly available transcripts, and four different transcription services were enlisted to transcribe these programs as part of the corpus creation effort (one for all TDT Mandarin data, one for TDT3 English, and two others for TDT2 English). The services varied in terms of the quality of transcripts delivered, with one of the TDT2 English services having been poorest.

In addition to manual transcription, all audio sources were submitted to unguided ASR, to establish a benchmark of TDT system performance given this quality of text as input. The TDT2 English data was submitted to two different English-based ASR systems; of course, in the absence of fully accurate manual transcriptions for most of this material, it remains difficult to compare their performance in terms of word error rate.

For all Mandarin sources, both newswire text and radio transcripts, additional data streams were produced by running the data through a Chinese-to-English machine translation system (SYSTRAN), with no manual guidance or revision of the system output. For audio data, both the manual and ASR transcriptions were translated in this way. Again, this was intended to set a benchmark for cross-lingual TDT performance given this quality of input.

As a result, most data sources were represented by at least two parallel, independent data streams – a couple of sources had three or four streams – each with its own peculiar properties and token sequence. The stable points of reference across all streams were the boundaries between news stories, and in the case of audio sources, the time offsets of those boundaries in the speech data.

3.2. Creating and tracking multiple annotations

A relational database was used to track the main stages of data creation and TDT-specific annotations. The basic units of the corpus are the sample file and the topical story unit. For audio sources, an entry was created each time a recording process was scheduled; the entry was updated at the conclusion of the recording, updated again after manual inspection to determine whether the recording was successful, and again after manual segmentation of the 30- or 60-minute file into story and non-news segments. This done, each story was assigned a unique identifier, which included the source, date, broadcast start time, and time offset within the file at the start of the story, and these identifiers were entered into the database to guide the topic annotation. (The equivalent stages for newswire data were fully automated

to prepare the stream or bulk archive text data for topic annotation.)

During the main topic annotation phase, in which every story had to be read some minimum number of times to assess it against all selected TDT topics, annotators also had to decide whether a given story was flawed in any of four ways, making it unsuitable for topic labeling. In the newswire data, a reported flaw would typically result in the removal of a story from the corpus, but for broadcast data, a flaw would generally be the result of a mistake during the manual story segmentation phase. Broadcast stories could not simply be discarded, as this would create gaps in the coverage of the continuous audio signal – segmentation errors needed to be fixed, and this would affect the inventory of news stories in the file, and/or the locations of boundaries (hence the identification numbers assigned to the stories would change, as well).

A further complication was the need in TDT3 to support alternative methods of topic annotation while the main annotation was still in progress. These alternative forms of labeling – first-story detection and story-link detection – were not actively tied into the database management of the main topic annotation; rather, they used a snap-shot of the corpus, taken as late and as carefully as possible during the main annotation. Fortunately, they would involve only a subset of the full TDT3 corpus, so it was possible to avoid particular files or stories where problems had been observed, and still provide an adequate sample. Despite our best efforts, some of the results of these alternate annotations could not be used in the final delivery of the corpus, because the stories to which they applied had been altered or removed in the course of repairs prompted by the main annotation; in particular, 0.5% of the 21,600 story link annotations were discarded because a number of stories had been eliminated from the corpus.

Additional uses of TDT data have already begun, spurring new annotations that were not part of the original corpus design. In order to establish a better estimate of ASR performance on this data set, NIST selected a random sample of 530 individual news stories from the TDT2 English corpus, totaling 10 hours of speech. The LDC adapted the text for these stories to the Hub-4 transcription specifications for broadcast news, and carefully went over each story, adding in the missing words and disfluencies, correcting spelling errors (common in closed-caption text), and adding time stamps to break long turns by news announcers into manageable phrases. This 10-hour set of careful transcription is now available from the LDC.

Other directions for annotation of TDT have included the identification of named entities, in support of the TREC project and related research, and the identification of new information across a sequence of stories on a given topic.

In a sense, each of these various annotations could be said to stand on its own as a sample for modeling a particular characteristic of language behavior or information flow. But research tasks moving along these various lines will tend to intersect, and it will be important to know to what extent their respective annotations intersect as well. The same crew may be handling the annotations for, e.g., first story detection, marking of new information, and named

entity extraction, but these tasks might not be carried out in unison (indeed they typically will not), they might be applied to discrete subsets of the corpus, and even if they do overlap on some sampling of the data, it might not be immediately obvious how to integrate these different annotations. Needless to say, it would not seem improbable that researchers could find the intersection of these annotations to be of some value.

4. Annotation Graphs as a Means of Integration

Annotation graphs were introduced by Bird and Liberman (Bird & Liberman, 1999a) as a convenient data model which abstracts away from the many contingent details of corpus file formats. An important consequence of this move, explored by Bird and Liberman (Bird & Liberman, 1999b), is that multiple independent annotations of a single corpus can be accessed and analyzed simultaneously. In this section we discuss the case of SWB, and the data shown in Figure 1.

Figure 2 shows the annotation graph for this SWB data, corresponding to the interval [21.86, 26.10]. In this graph, word arcs have type W/, Treebank arcs have T/ and disfluency arcs have DISF/ type. Types for the part-of-speech arcs have been omitted for sake of clarity (i.e. Pos/metric/JJ is written as just metric/JJ). For readability, the graph is represented in two pieces; the lower piece should be interpolated into the upper piece at the position of the dotted arc labeled X. Observe that the equivocation about the tokenization of punctation from Figure 1 is preserved in the annotation graph.

The ability to merge diverse layers of independent annotations into a single graph derives from the definition of nodes for connecting the arcs. The nodes are anchored at specific points along a time line representing the signal being annotated. In SWB, the LDC interim transcripts, the ICSI phonetic transcripts, and the ISIP revised transcripts each represent a distinct segmentation of the time lines for the corresponding signals. The migration of these data sets into annotation graphs provides a stable, time-anchored basis for cross-reference; owing to the divergent turn- or phrase-level segmentations provided by each data set, the time line is in fact the only practical basis for cross-reference.

Even though the phonetic transcripts cover only sparse portions of most calls, these partial annotations for any one file can be mapped coherently onto a single complete graph that includes the other transcripts in their entirety, providing a well-defined algorithmic approach for locating and resolving discrepancies in lexical content, phrasal segmentation, and word-level time alignments. The disfluency, discourse and Treebank annotations are derived from the LDC transcripts, and as demonstrated above, they can also be incorporated into the one graph, making it quite simple to identify the particular elements in those annotations that will be affected by corrections to the underlying transcripts. It follows that the propagation of transcription repairs through all levels of annotation becomes a well-behaved and accountable process.

Apart from the obvious benefits to corpus maintenance, this approach to handling annotations provides an important capability for integrating the results of diverse annotation efforts in new research. If and when a prosodic annotation of SWB becomes available, we can readily envision the ability to study interactions among intonational focus, discourse function, syntactic structure, and phonological processes, simply by adding the one new layer to the existing network of other annotations.

The application of annotation graphs to TDT is equally fruitful. The stability of reference to the basic units of TDT corpora (sample files, news stories, word tokens) is already well established, owing to the fact that the corpus creation effort has been tightly centralized. But due to the overall bulk of the data, most new annotations, especially those requiring human judgment, are likely to be limited to cover only portions of the collection. Again, these partial annotations can form coherent annotation graphs on their own, and can be treated atomically or integrated with other graph structures as needed.

The issue of data formatting for creation, storage, distribution and research use of annotations is an independent concern, orthogonal to the use of annotation graphs as a framework for handling corpora. The arc-and-node structure can be rendered into (and retrieved from) a very simple XML data stream, and it is equally possible to create filters that can populate an annotation graph by reading any chosen data format, without loss of information. Filters could also be made to create a chosen data format from an annotation graph, though it's possible that some information in the graph would not be preserved in the process.

5. Conclusion

We have presented an overview of two large speech corpora, both of which have received a wide range of divergent and independent annotations. For Switchboard, we have discussed some details about the comparability and compatibility of the various annotations, and have presented an analysis framework that will enable a high degree of integration among them, in terms of both maintenance and research use. The case of TDT demonstrates that even with a centralized corpus creation effort, there can still be problems with handling data repairs and consistency when distinct sets of annotation must be carried out simultaneously. Both corpora present the need to accommodate sparse annotations in a manner that does not sacrifice the overall coherence of the larger corpus. Annotation graphs provide an effective framework for meeting this need.

6. References

Bird, S. and M. Liberman, 1999a. A Formal Framework for Linguistic Annotation. *Technical Report MS-CIS-99-01*, Department of Computer and Information Science, University of Pennsylvania. (Available at *http://morph.ldc.upenn.edu/Papers/*).

Bird, S. and M. Liberman, 1999b. Annotation graphs as a framework for multidimensional linguistic data analysis. *Towards Standards and Tools for Discourse Tagging – Proceedings of the Workshop*, Association for
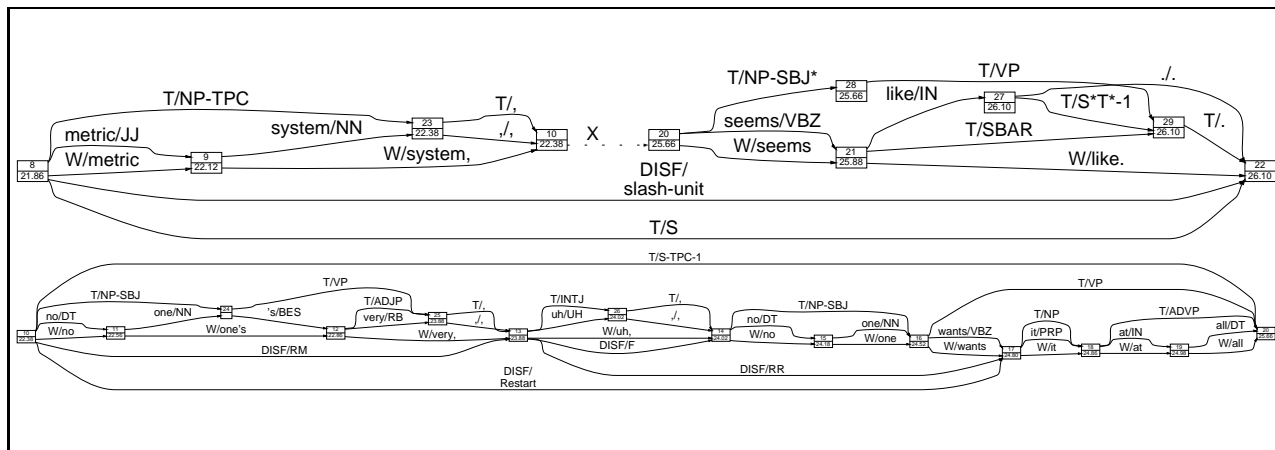
Figure 2: Annotation Graph Fragment for Switchboard Data

Computational Linguistics. Somerset, NJ. (Available at *http://morph.ldc.upenn.edu/Papers/*).

Cieri, C., D. Graff, M. Liberman, N. Martey, S. Strassel, 2000. The TDT-2 Text and Speech Corpus. *Proceedings of LREC-2000*.

The Center for Language and Speech Processing. FTP server for Summer Workshops at Johns Hopkins University. *http://www.clsp.jhu.edu/*.

Godfrey, J., E. Holliman and J. McDaniel, 1992. SWITCH-BOARD: Telephone speech corpus for research and development. *Proceedings of ICASSP-92*, pp. 517–520. San Francisco.

Greenberg, S., 1999. Switchboard Transcription Project. *http://www.icsi.berkeley.edu/real/stp/*.

Jurafsky, D., 1997. Web page for the 1997 CLSP/JHU Summer Workshop. *http://www.Colorado.EDU/linguistics/jurafsky/ws97/*.

Jurafsky, D., L. Shriberg, D. Biasca, 1997. Switchboard SWBD-DAMSL Shallow-Discourse-Function Annotation. *http://stripe.colorado.edu/˜jurafsky/manual.august1.html*.

Mantha, V. , 2000. Switchboard Resegmentation Project. *http://www.isip.msstate.edu/projects/switchboard/*.

Wayne, C. L., 2000. Multilingual Topic Detection and Tracking: Successful Research Enabled by Corpora and Evaluation. *Proceedings of LREC-2000*.